

Exp No: 10

Date: 15.09.2024

**IMPLEMENT THE MAX TEMPERATURE MAPREDUCE PROGRAM TO
IDENTIFY THE YEAR WISE MAXIMUM TEMPERATURE FROM SENSOR
DATA**

AIM:

To implement the Max temperature Map Reduce program to identify the year-wise maximum temperature from the sensor data.

PROCEDURE:

1. Open command prompt and run as administrator

Go to hadoop sbin directory

```
C:\Windows\system32>cd C:\Hadoop\sbin  
C:\Hadoop\sbin>_
```

Note:

1. Check hadoop/data/datanode and hadoop/data/namenode and if both folders are empty, type “hdfs namenode -format”.
2. Check python version with “python --version”.
3. Check “C:\Python39\” is added in Environment variables > System variables > Path, if not add your python path.
4. Check Environment variables > System variables > HADOOP_HOME is set as “C:\Hadoop”.

2. Start Hadoop Services

start-dfs.cmd

start-yarn.cmd

```
C:\Hadoop\sbin>start-dfs.cmd  
  
C:\Hadoop\sbin>start-yarn.cmd  
starting yarn daemons  
  
C:\Hadoop\sbin>jps  
13120 NameNode  
2384 NodeManager  
4100 DataNode  
7956 ResourceManager  
9124 Jps
```

Hadoop

Overview

Datanodes

Datanode Volume Failures

Snapshot

Startup Progress

Utilities

Overview 'localhost:9000' (✓active)

Started:	Sun Aug 18 18:45:16 +0530 2024
Version:	3.3.6, r1be78238728da9266a4f88195058f06fd012bf9c
Compiled:	Sun Jun 18 13:52:00 +0530 2023 by ubuntu from (HEAD detached at release-3.3.6-RC1)
Cluster ID:	CID-a23ce25d-ee9d-4000-ac1f-044f436c4c8a
Block Pool ID:	BP-934656018-192.168.56.1-1723971050909

Summary

Security is off.

Safemode is off.

19 files and directories, 5 blocks (5 replicated blocks, 0 erasure coded block groups) = 24 total filesystem object(s).

Heap Memory used 74.86 MB of 193 MB Heap Memory. Max Heap Memory is 889 MB.

Non Heap Memory used 61.65 MB of 63.11 MB Committed Non Heap Memory. Max Non Heap Memory is <unbounded>.

Configured Capacity:	118.63 GB
----------------------	-----------

```
hdfs dfs -mkdir -p /maxweather/hadoop/input
```

```
C:\Hadoop\sbin>hdfs dfs -mkdir -p /maxweather/hadoop/input
```

```
hdfs dfs -put C:/Users/Admin/max_weather_cc/weather_data.txt /maxweather/hadoop/input/
```

```
C:\Hadoop\sbin>hdfs dfs -put C:/Users/Admin/max_weather_cc/weather_data.txt /maxweather/hadoop/input/

C:\Hadoop\sbin>hdfs dfs -ls /maxweather/hadoop/input/
Found 1 items
-rw-r--r--  1 Admin supergroup          3264 2024-09-15 14:49 /maxweather/hadoop/input/weather_data.txt

C:\Hadoop\sbin>hdfs dfs -cat /maxweather/hadoop/input/weather_data.txt
0029029070999992021010720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00311+99999098351ADDFG102991999999999999999999
002902907099999202108720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00121+99999098351ADDFG102991999999999999999999
0029029070999992021030720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00231+99999098351ADDFG102991999999999999999999
0029029070999992021040720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00411+99999098351ADDFG102991999999999999999999
0029029070999992021050720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00501+99999098351ADDFG102991999999999999999999
0029029070999992021060720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00211+99999098351ADDFG102991999999999999999999
00290290709999920210720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00221+99999098351ADDFG102991999999999999999999
002902907099999202208720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00311+99999098351ADDFG102991999999999999999999
0029029070999992022030720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00401+99999098351ADDFG102991999999999999999999
0029029070999992022040720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00501+99999098351ADDFG102991999999999999999999
0029029070999992022050720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00311+99999098351ADDFG102991999999999999999999
0029029070999992022060720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00241+99999098351ADDFG102991999999999999999999
0029029070999992023010720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00301+99999098351ADDFG102991999999999999999999
002902907099999202308720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00421+99999098351ADDFG102991999999999999999999
0029029070999992023030720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00531+99999098351ADDFG102991999999999999999999
0029029070999992023040720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00211+99999098351ADDFG102991999999999999999999
0029029070999992023050720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00301+99999098351ADDFG102991999999999999999999
0029029070999992023060720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00411+99999098351ADDFG102991999999999999999999
0029029070999992024010720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00111+99999098351ADDFG102991999999999999999999
0029029070999992024020720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00231+99999098351ADDFG102991999999999999999999
0029029070999992024030720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00301+99999098351ADDFG102991999999999999999999
0029029070999992024040720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00411+99999098351ADDFG102991999999999999999999
0029029070999992024050720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00501+99999098351ADDFG102991999999999999999999
0029029070999992024060720004+64333+023450FM-12+000599999V0202501N027819999999N0000001N9+00211+99999098351ADDFG102991999999999999999999
```

Note:**mapper.py:**

```
#!/usr/bin/env python3
import sys
for line in sys.stdin:
    line = line.strip()
    if not line:
        continue
    try:
        # Extract relevant fields from the raw data
        year = line[15:19] # Extract year from the line
        temp_str = line[90:92] # Extract temperature from the line
        quality = line[92:93] # Extract quality indicator

        # Check if the temperature is valid and the quality is acceptable
        if temp_str != "+9999" and quality in ['0', '1', '4', '5', '9']:
            temp = int(temp_str)
            print(f"{year}\t{temp}")
    except Exception as e:
        sys.stderr.write(f"Error processing line: {line}\nException: {str(e)}\n")
```

combiner.py:

```
#!/usr/bin/env python
import sys
from collections import defaultdict
current_year = None
temp_set = set()
for line in sys.stdin:
    line = line.strip()
    if not line:
        continue
    try:
        year, temp_str = line.split('\t')
        temp = int(temp_str)
        if current_year == year:
            temp_set.add(temp)
```

```

else:
    if current_year:
        # Print the year and the set of temperatures
        for t in temp_set:
            print(f"{current_year}\t{t}")
        current_year = year
        temp_set = {temp}
except Exception as e:
    sys.stderr.write(f"Error processing line: {line}\nException: {str(e)}\n")

# Output the set of temperatures for the last year
if current_year:
    for t in temp_set:
        print(f"{current_year}\t{t}")

reducer.py:
#!/usr/bin/env python
import sys
current_year = None
max_temp = None
for line in sys.stdin:
    line = line.strip()
    if not line:
        continue
    try:
        year, temp_str = line.split("\t")
        temp = int(temp_str)
        if current_year == year:
            if temp > max_temp:
                max_temp = temp
        else:
            if current_year:
                # Print the maximum temperature for the previous year
                print(f"{current_year}\t{max_temp}")
            current_year = year
            max_temp = temp
    except Exception as e:
        sys.stderr.write(f"Error processing line: {line}\nException: {str(e)}\n")

```

```
# Output the maximum temperature for the last year
```

```
if current_year:
```

```
    print(f"{current_year}\t{max_temp}")
```

6. Run the Hadoop Streaming Job

```
hadoop jar %HADOOP_HOME%\share\hadoop\tools\lib\hadoop-streaming-*.jar ^
```

```
-input /maxweather/hadoop/input/weather_data.txt ^
```

```
-output /maxweather/hadoop/output ^
```

```
-mapper "python C:\Users\Admin\max_weather_cc\mapper.py" ^
```

```
-combiner "python C:\Users\Admin\max_weather_cc\combiner.py" ^
```

```
-reducer "python C:\Users\Admin\max_weather_cc\reducer.py" ^
```

```
-file C:\Users\Admin\max_weather_cc\mapper.py ^
```

```
-file C:\Users\Admin\max_weather_cc\combiner.py ^
```

```
-file C:\Users\Admin\max_weather_cc\reducer.py
```

```
C:\Hadoop\sbin>hadoop jar %HADOOP_HOME%\share\hadoop\tools\lib\hadoop-streaming-*.jar ^
More? -input /maxweather/hadoop/input/weather_data.txt ^
More? -output /maxweather/hadoop/output ^
More? -mapper "python C:\Users\Admin\max_weather_cc\mapper.py" ^
More? -combiner "python C:\Users\Admin\max_weather_cc\combiner.py" ^
More? -reducer "python C:\Users\Admin\max_weather_cc\reducer.py" ^
More? -file C:\Users\Admin\max_weather_cc\mapper.py ^
More? -file C:\Users\Admin\max_weather_cc\combiner.py ^
More? -file C:\Users\Admin\max_weather_cc\reducer.py
2024-09-15 16:05:06,874 WARN streaming.StreamJob: -file option is deprecated, please use generic option
packageJobJar: [C:\Users\Admin\max_weather_cc\mapper.py, C:\Users\Admin\max_weather_cc\combiner.py, C:\U
[] C:\Users\Admin\AppData\Local\Temp\streamjob4190031441542396910.jar tmpDir=null
2024-09-15 16:05:09,340 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at
2024-09-15 16:05:09,758 INFO client.DefaultNoHARMFailoverProxyProvider: Connecting to ResourceManager at
2024-09-15 16:05:11,009 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadd
2024-09-15 16:05:12,204 INFO mapred.FileInputFormat: Total input files to process : 1
2024-09-15 16:05:12,612 INFO mapreduce.JobSubmitter: number of splits:2
2024-09-15 16:05:12,963 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1726391798653_0008
2024-09-15 16:05:12,963 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-09-15 16:05:13,260 INFO conf.Configuration: resource-types.xml not found
2024-09-15 16:05:13,262 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-09-15 16:05:13,471 INFO impl.YarnClientImpl: Submitted application application_1726391798653_0008
2024-09-15 16:05:13,575 INFO mapreduce.Job: The url to track the job: http://DESKTOP-TF65P79:8088/proxy,
2024-09-15 16:05:13,583 INFO mapreduce.Job: Running job: job_1726391798653_0008
2024-09-15 16:05:33,469 INFO mapreduce.Job: Job job_1726391798653_0008 running in uber mode : false
2024-09-15 16:05:33,472 INFO mapreduce.Job: map 0% reduce 0%
2024-09-15 16:05:54,357 INFO mapreduce.Job: map 100% reduce 0%
2024-09-15 16:06:06,579 INFO mapreduce.Job: map 100% reduce 100%
2024-09-15 16:06:08,617 INFO mapreduce.Job: Job job_1726391798653_0008 completed successfully
2024-09-15 16:06:08,873 INFO mapreduce.Job: Counters: 55
```

```
File Input Format Counters
```

```
Bytes Read=4896
```

```
File Output Format Counters
```

```
Bytes Written=32
```

```
2024-09-15 16:06:08,910 INFO streaming.StreamJob: Output directory: /maxweather/hadoop/output
```

7. View the Output

`hdfs dfs -cat /maxweather/hadoop/output/part-00000`

```
C:\Hadoop\sbin>hdfs dfs -cat /maxweather/hadoop/output/part-00000
2021    50
2022    50
2023    53
2024    50
```

8. Once the map reduce operations are performed successfully, the output will be present in the specified directory.

“/maxweather/hadoop/output/part-00000”

Browse Directory

Show 25 entries

Search:

<input type="checkbox"/>	Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name	
<input type="checkbox"/>	-rw-r--r--	Admin	supergroup	0 B	Sep 15 15:34	1	128 MB	_SUCCESS	
<input type="checkbox"/>	-rw-r--r--	Admin	supergroup	0 B	Sep 15 15:34	1	128 MB	part-00000	

Showing 1 to 2 of 2 entries

Previous

1

Next

File contents

2021 50

2022 50

2023 53

2024 50

9. Stop Hadoop Services

`stop-dfs.cmd`

`stop-yarn.cmd`

```
C:\Hadoop\sbin>stop-dfs.cmd
SUCCESS: Sent termination signal to the process with PID 3636.
SUCCESS: Sent termination signal to the process with PID 14040.

C:\Hadoop\sbin>stop-yarn.cmd
stopping yarn daemons
SUCCESS: Sent termination signal to the process with PID 10644.
SUCCESS: Sent termination signal to the process with PID 8316.

INFO: No tasks running with the specified criteria.

C:\Hadoop\sbin>
```

RESULT:

Thus, to implement the Max temperature Map Reduce program to identify the year-wise maximum temperature from the sensor data was completed successfully.